



BRAND PREFERENCE PREDICTION

Ubiquum – DA121- Technical department

Alessandro Arnone

26/04/20

TABLE OF CONTENTS

OVERVIEW	2
the data	2
ANALYSIS	4
select the right variables	4
select the right model	6
model tuning	6
C5.0 Automatic Grid.....	6
C5.0 Manual Grid.....	7
Random Forest - Automatic Grid	7
K-nn – Automatic Grid:	7
SvmRadial – Automatic Grid	8
SvmLinear	8
model selection	8
CONCLUSION	9

BRAND PREFERENCE PREDICTION REPORT

OVERVIEW

The objective of this report is to investigate whether customer responses to survey questions allow us to predict the computer brand preference and if so, to provide the sales team with a complete view of what brand our customers prefer based on the predictions made.

We will run and optimize different decision tree classification methods in order to find patterns in our dataset.

The final model chosen to predict the uncompleted surveys is svm R which produced the following results:

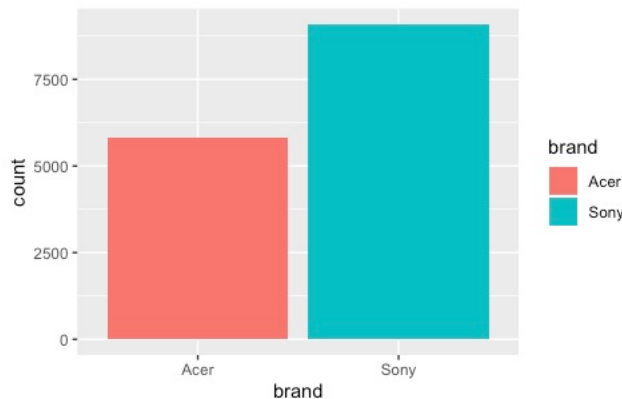


Figure 1- brand preference

- Around 65% of the people prefer Sony
- There is a relation between Brand, Salary and Age

THE DATA

The sales team engaged a market research firm to conduct a survey of our existing customers, providing us with all the information gathered in a CSV file.

Specifically, as first insights we can conclude that:

- 9898 surveys have been submitted
- 7 different variables have been recorded through the survey (Salary, Age, Eleven, Car, Zipcode, Credit, Brand)
- Around 65% of the people involved in the study prefer a Sony computer. This class difference could potentially create some inbalance in our prediction (since a model created on inbalanced classes tends to predict for the most numerous class) hence this needs to be treated. As solution, the dataset will be down-sampled (in this case, we preferred to lose information rather than having an additional 40% of information that can be Bias) to a total of 7488 observations (50% of customers preferring Acer, 50% of customers preferring Sony – see Figure 2 for the distribution before balancing the class)
- All the variables are uniformly distributed (see tab.1 for histograms and plots)
- Data looks either manipulated and altered (credit and salary have decimals, which is uncommon) or automatically generated by a machine
- No outliers identified
- No rows duplicates found

BRAND PREFERENCE PREDICTION REPORT

Table 1- Variables Distribution


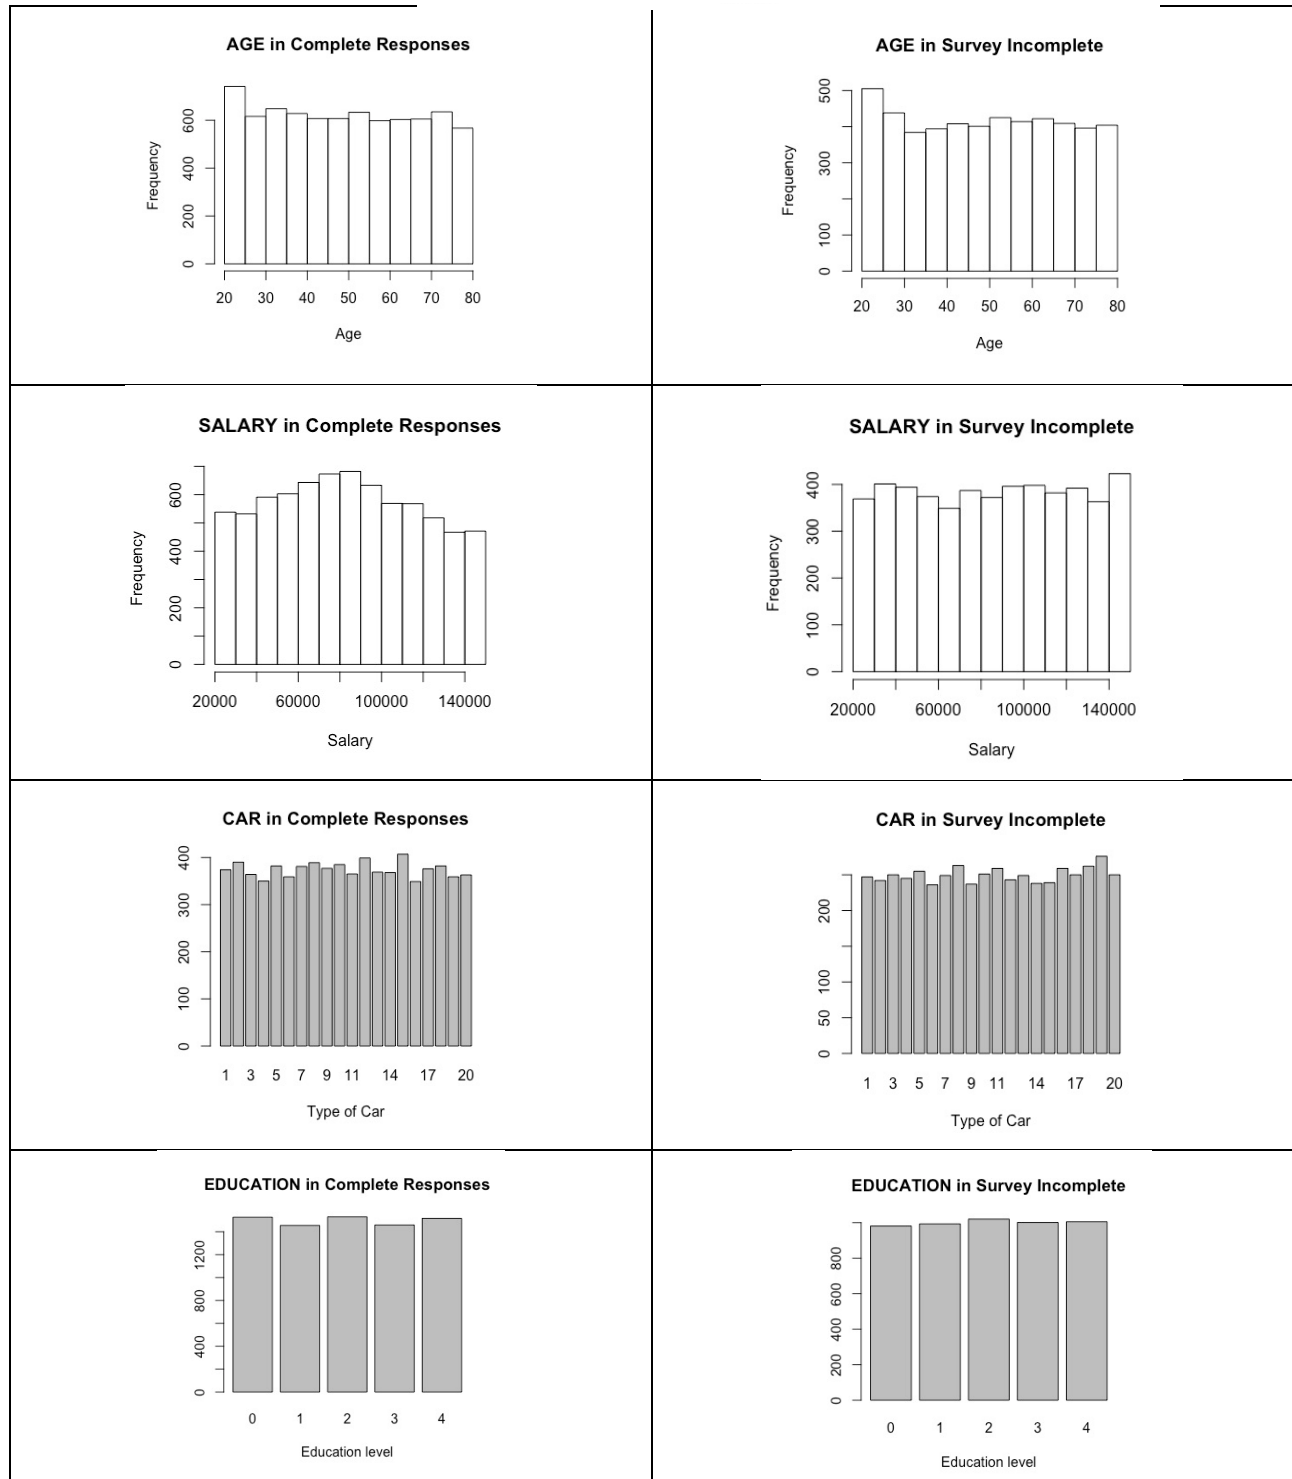



Figure 2 - Inbalanced classes



BRAND PREFERENCE PREDICTION REPORT



ANALYSIS

SELECT THE RIGHT VARIABLES

In order to predict the Brand Preference, the variables that most influence this factor need to be identified. For the selection of predictors we will use both common sense and statistic tools such as regression analysis, the function `varimp()`, statistical hypothesis tests (such as Chi-squared test), logistic regression and anova.

The variable to exclude are all the variables that:

1. logically are not linked with the Brand Preference such as **Car**
2. have p-value > 0.99 during statistical hypothesis tests (Chi-squared test) and low Cramer's V such as **Education level** (see Tab. 2)
3. have no statistical relevance when modelling the variables with the logistic regression such as **Credit** (see Tab.3)
4. do not contribute heavily to the model during the estimation of the contribution of each variable with the `var.imp()` function such as **ZIPCODE** (see Figure 3)

Table 2 - Chi-squared test and Cramer's V

	p-value	Cramer's V	Comments	
Level of education	0.99	0.006	p-value almost 1. Since our condition of dependency is $p < 0.05$, such a big p with such a low Cramer can definitely suggest us that the 2 variables are independent.	Independent
Car	0.52	0.049	p-value not small enough to reject the hypothesis of independency and Cramer's V really low. Brand and car are independent	Independent
Zip	0.19	0.03	p-value not low enough to reject the hypothesis of independency but not big enough to define hypothesis of dependencies	Need additional information

Table 3 - Logistic regression


	Residual Deviance	p-value	Comments	
Null (only intercept)	10380			
Salary	9965	< 2e-16	p-value really small. It suggests that there is a statistical significance and we can reject the hypothesis of independence. More over we see a drop in the residual deviance compared to the Null model which is a good parameter to understand if the variable is contributing to the model	Dependent

BRAND PREFERENCE PREDICTION REPORT

Age	9960	0.02454	p-value really small. It suggests that there is a statistical significance and we can reject the hypothesis of independence. It contributes less to the model	Dependent
Credit	9960	0.30852	p-value big. There is no statistical significance and we cannot reject the hypothesis of independencies. Also it does not contribute to the model since the Residual Deviance keep on being the same	Independent

In fact, an additional check can be done through the `varimp()` function of `caret` which evaluate the importance of a variable using a model-based approach. In our case we used this function when trying to apply a random forest model to our result. The data can be sum-up in the graph below where Salary and Age are the one with the biggest impact on the prediction compared to Car, Credit, Zipcode and elevel.

Figure 3 - Plot `varimp()`



To sum up the variables picked as predictors are Salary and Age. A first prediction of the model can be seen plotting the decision tree where we can see that the only two variable shown are the one mentioned earlier.




Figure 4 - Decision Tree

BRAND PREFERENCE PREDICTION REPORT

SELECT THE RIGHT MODEL

Once selected the variables to include in our predictive model, multiple tests will be run in order to find the model that maximize the different accuracy metrics used to assess the model such as: **Accuracy** and **Kappa**

Specifically, the predictive Model used are: C5.0, Random Forest, K-nn, SVM Radial and SVM linear.

The following steps have been run in order to perform this task:

1. **Model tuning:** Tune each of the model with different hyperparameters and find the ones that maximize the accuracies metrics
2. **Model selection:** Choose the best model

MODEL TUNING

A cross-validation 10-folder has been applied to all the models. The best parameters for each model are underlined in yellow.

C5.0 AUTOMATIC GRID

- TuneLength = 3
- PreProcess= Center, Scale

Table 4 - C5.0 accuracy table

model	winnow	trials	Accuracy	Kappa
rules	FALSE	1	0.839	0.678
rules	FALSE	10	0.918	0.836
rules	FALSE	20	0.920	0.840
rules	TRUE	1	0.839	0.678
rules	TRUE	10	0.918	0.836
rules	TRUE	20	0.920	0.840
tree	FALSE	1	0.839	0.678
tree	FALSE	10	0.919	0.838
tree	FALSE	20	0.920	0.839
tree	TRUE	1	0.839	0.678
tree	TRUE	10	0.919	0.838
tree	TRUE	20	0.920	0.839

Additional Consideration:

- No rules have been defined
- Winnow results having the same value since the algorithm uses only 2 variables (the ones not important for the model have already been excluded)
- There in an increase of accuracy increasing the number of trials (as we can see from the graph below)

BRAND PREFERENCE PREDICTION REPORT

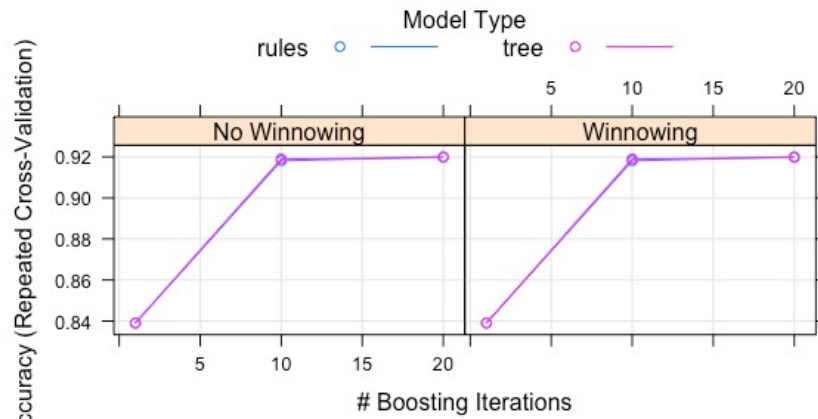


Figure 5 - accuracy plot

C5.0 MANUAL GRID

- Winnow= False
- Model= Tree
- Trials= 25, 35, 45

Table 5 - C5.0 Accuracy Table

Trials	Accuracy	Kappa
25	0.922	0.844
35	0.922	0.843
45	0.922	0.843

Additional Consideration:

- The increase of boosting iteration does not bring any increase in the accuracy

RANDOM FOREST - AUTOMATIC GRID

- Tune Length=3
- Mtry= 2,4,6
- For academic purpose the variable excluded has been reintroduced to study the effect of different variables on the models

Table 6 - Accuracy Table

mtry	Accuracy	Kappa
2	0.914	0.828
4	0.922	0.845
6	0.918	0.835

Additional Consideration:

- there is an accuracy improvement if we include other 2 variables to the model. This could potentially create problem since there might be a risk of overfitting the model and an increase of computational time with only a 0.08% gain

K-NN – AUTOMATIC GRID:

BRAND PREFERENCE PREDICTION REPORT

- TuneLength= 5
- PreProcess= Center, Scale
- Value k = 5,7,9,11,13
- Computational time is really low

Table 7- K-NN Accuracy Table

k	Accuracy	Kappa
5	0.919	0.837
7	0.920	0.840
9	0.922	0.844
11	0.923	0.847
13	0.923	0.845

SVMRADIAL – AUTOMATIC GRID

- TuneLength= 5
- PreProcess= Center, Scale
- Value C = {0.25 , 0.50 , 1.00 , 2.00 , 4.00 }
- Gamma = 1.25

Table 8 - Accuracy Table SVM Radial

C	Accuracy	Kappa
0.25	0.920	0.839
0.50	0.923	0.847
1.00	0.924	0.848
2.00	0.925	0.849
4.00	0.926	0.853

SVMLINEAR

- TuneLength= 5
- PreProcess= Center, Scale
- Tuning parameter 'C' was held constant at a value of 1

Table 9 - Accuracy Table SVM Linear

C	Accuracy	Kappa
1	0.682	0.363

MODEL SELECTION

The best tuning parameters are summarized in the following tab:

Model	Parameters	Accuracy	K
C5.0	Trials=20	0.920	0.839
C5.0 Manual	Trials=25	0.922	0.844
Random Forest	mtry= 4	0.922	0.845

BRAND PREFERENCE PREDICTION REPORT

Random Forest	mtry=2	0.921	0.842
K-nn	11	0.923	0.847
SVM Linear	C=1	0.682	0.363
SVM radial	C=4 Sigma = 1.23	0.926	0.853

The model to exclude are:

- SVM linear since It presents the lower accuracy
- Random Forest with mtry= 4 since the computational time and the risk of overfitting do not justify the choice

Moreover, if we check the accuracy of our prediction on our Test Set:

Model	Parameters	Accuracy	K
C5.0	Trials=20	0.923	0.845
Random Forest	mtry=2	0.922	0.844
kNN	N=11	0.928	0.857
svmRadial	C=4 sigma = 1.23	0.929	0.858

- svmRadial present the higher level of accuracy and k
- the 4 models only differentiates by less then 7 centesimal. Hence the decision of a model simply based on a variable like Accuracy is risky. That said, we will continue our analysis with svmRadial

The confusion matrix of the model choosed applied to the test Dataset is :

Table 10 - Confusion matrix

	Prediction	Acer	Sony	
Observed				
	Acer	893	47	91%
	Sony	85	847	94%
		91%	94%	

A high % of each class justify the high K accuracy we found in the model. Moreover, since the classes have been balanced, the number of Acer and the number of Sony is pretty similar.

CONCLUSION

Applying the model choosen (svm Radial) to the uncompleted surveys we have the following distribution:





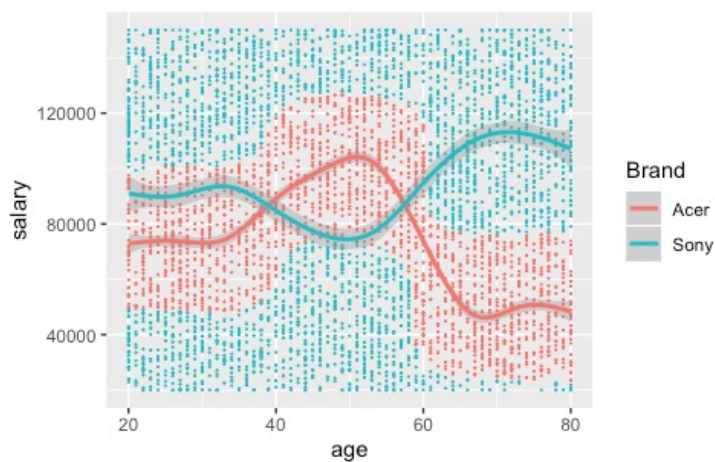
Figure 7- Brand prediction

Figure 6 - Unbalanced Class in the original Dataset

BRAND PREFERENCE PREDICTION REPORT



We have 2072 people preferring Acer and 2928 people preferring Sony. Exactly the same proportion of the dataset before the class balance. Moreover, as we can see from the graphes below:



not only we have the same behaviour in both dataset analyzed (Complete Response before balanced, Incomplete Survey) but also a similiar proportion of 65/35 of people preferring either Sony or Acer respectively. That means:

- Dataset has been altered to keep the proportion 65%-35%
- It seems that the Survey Incomplete dataset it's partially a subset of the Complete Survey Dataset
 - some salaries are the same (included the decimals)
 - if we round up salaries to the closest integer, we find around 1800 rows identical in terms of salaries and age
- Whilst the prediction shoulds reliable based on the model performances, we need to take into account the fact that the reliability depends on the quality of the dataset